

R Programming week 2

Basic Constructs

- if, else
- for
- while
- repeat
- break - break execute of a loop
- next - skip an iteration of a loop
- return - exit a function

```
# if
if(<condition>) {

} else {

}

if(<condition>) {

} else if(<condition>) {

} else {

}
```

if example

```
# If format: 1
x <- 3
if( x > 3 ) {
  y <- 10
} else {
  y <- 100
}
print(y)
```

```
# If format: 2 (r specific)
```

```
y <- if( x > 3 ) {  
  10  
} else {  
  100  
}  
print(y)
```

```
# FOR loop
```

```
for( x in 1:10) {  
  print(x)  
}
```

```
x <- c("a", "b", "c", "d")
```

```
for( i in 1:4) {  
  print(x[i])  
}
```

```
# by using seq_along() you do not need to know the length in advance
```

```
for( i in seq_along(x)) {  
  print(x[i])  
}
```

```
for( i in 1:4) print(x[i])
```

```
#
```

```
# nexted for-loops
```

```
#
```

```
x <- matrix(1:6, 2, 3)
```

```
for( i in seq_len(nrow(x))) {  
  for( j in seq_len(ncol(x))) {  
    print(x[i, j])  
  }  
}
```

```
##
```

```
## WHILE LOOP
```

```
##
```

```
# loop from 0 to 9
```

```
count <- 0
```

```
while(count < 10 ) {
  print(count)
  count <- count +1
}
```

test multiple conditions in a while loops

```
z <- 5
```

```
while( z >= 3 && z <= 10) {
  print(z)
  coin <- rbinum(1, 1, 0.5)

  if(coun == 1 ) { ## random walk
    z <- z + 1
  } else {
    z <= z -1
  }
}
```

-- test 2 - multiple expression in a condition
-- are evaluated left to right

Test

```
underflowCount <- 0
overflowCount <- 0
for( i in 1:10000) {
  z <- 7
```

```
    while( z >= 3 && z <= 11) {
      print(z)
      coin <- rbinom(1, 1, 0.5)

      if(coin == 1 ) { ## random walk
        z <- z + 1
      } else {
        z <- z -1
      }
    }
    if( z < 3 ) {
      print("z underflow")
      underflowCount <- underflowCount + 1
    } else {
      print("z overflow")
      overflowCount <- overflowCount + 1
    }
  }
```

```
}  
print(paste("Overflows", overflowCount, "Underflows", underflowCount, "Balance %",  
(underflowCount-overflowCount)/100))
```

```
##
```

```
## repeat
```

```
##
```

```
x0 <- 1
```

```
tol <- 1e-8
```

```
repeat {
```

```
  x1 <- computeEstimate()
```

```
  if(abs(x1 - x0) < tol) {
```

```
    break
```

```
  } else {
```

```
    x0 <- x1
```

```
  }
```

```
}
```

```
# no guarantee when it ends
```

```
# next, return, break
```

```
for( i in 1:100) {
```

```
  if( i <= 20) {
```

```
    ## Skip the first 20 iterations
```

```
    next
```

```
  }
```

```
  ## Do something here
```

```
}
```

```
# return signals that a function should exit and return a given value
```

```
##
```

```
## FUNCTIONS
```

```
##
```

```
## In the future you should put your r functions into a r package
```

```
##
```

```
# very simple function
```

```
add2 <- function(x, y) {
```

```
  x + y
```

```
}  
add2(10,14)  
# [1] 24
```

```
# above 10  
above10 <- function(x) {  
  # x[x > 10]  
  # or use a variable  
  use <- x > 10  
  x[use]  
}
```

```
above10(c(1:20))
```

```
# above X default 10  
above <- function(d,x=10) {  
  # d[d > x]  
  # or use a variable  
  use <- d > x  
  d[use]  
}
```

```
above(c(1:20), 19)  
above(c(1:20))
```

more complex - calculate mean of a matrix or dataframe

```
columnmean <- function(y, removeNA = TRUE) {  
  nc <- ncol(y)  
  means <- numeric(nc) # initial vector of columns (to zero)  
  for( i in 1:nc) {  
    means[i] <- mean(y[,i], na.rm = removeNA)  
  }  
  means  
}
```

```
columnmean(airquality)  
columnmean(airquality, FALSE)  
columnmean(airquality, TRUE)
```

```
# R functions are R objects like other types (type "function" in this case)  
class(columnmean)  
# [1] "function"
```

functions are first class objects like character, numeric etc.

Function Arguments

functions have arguments which potentially have default values.

- The formal arguments are the arguments included in the function definition

- The formals function returns a list of all the formal arguments of a function

- Not every function call in R makes use of all the formal arguments

- Function arguments can be missing or might have default values

Get the formal arguments

```
formals(columnmean)
```

```
# $y
```

```
#
```

```
# $removeNA
```

```
# [1] TRUE
```

Argument Matching

using the sd() standard deviation as an example

```
mydata <- rnorm(100)
```

```
sd(mydata)
```

```
# [1] 1.026745
```

```
sd(x = mydata, na.rm = FALSE)
```

```
# [1] 1.026745
```

```
sd(na.rm = FALSE, x = mydata)
```

```
# [1] 1.026745
```

```
sd(na.rm = FALSE, mydata)
```

```
# [1] 1.026745
```

Note: it's not recommended to change the order but you can

```
args(lm)
```

```
# function (formula, data, subset, weights, na.action, method = "qr",
```

```
#   model = TRUE, x = FALSE, y = FALSE, qr = TRUE, singular.ok = TRUE,
```

```
# contrasts = NULL, offset, ...)
```

```
# The following two calls are equivalent
```

```
# lm(data = mydata, y ~ x, model = FALSE, 1:100)
```

```
# lm(y ~ x, mydata, 1:100, model = FALSE)
```

```
# Defing a function
```

```
f <- function(a, b = 1, c = 2, d = NULL) {
```

```
}
```

```
# In addition to not specifying a default value, you can also
```

```
# also set an argument to NULL.
```

```
## Lazy Evaluation
```

```
f <- function(a, b ) {
```

```
  a^2
```

```
}
```

```
f(2)
```

```
# [1] 4
```

```
# Note: No Error in not passing (b) as (b) was not used in the function
```

```
f <- function(a, b ) {
```

```
  print(a)
```

```
  print(b)
```

```
}
```

```
f(45)
```

```
# where b is missing and used in the function - you will get an error.
```

```
f <- function(a, b ) {
```

```
  print(a)
```

```
  print(b)
```

```
}
```

```
f(45) # <-- will throw an error as b was not passed but is evaluated
```

```
[1] 45 # <-- this is the print(a) part - ok no error so far
```

```
# Error in print(b) : argument "b" is missing, with no default
```

```
##
## * VERY IMPORTANT **
##
## The (special) "..." Argument
##
## in this example, may a custom plot function
##
mplot <- function(x, y, type = "l", ...) {
  plot(x, y, type = type, ...)
}

# Generic functions use '...' so that extra arguments can be passed
# to the methods (more on this later).

mean
# function (x, ...)
# UseMethod("mean")

## the '...' is used in function like paste and cat
args(paste)
# function (..., sep = " ", collapse = NULL, recycle0 = FALSE)

args(cat)
# function (..., file = "", sep = " ", fill = FALSE, labels = NULL,
#   append = FALSE)

# all arguments after the '...' must be named, you also can't use partial Matching

args(paste)
# function (..., sep = " ", collapse = NULL, recycle0 = FALSE)

paste("a", "b", sep = ":")
# [1] "a:b"

paste("a", "b", se = ":") # <-- treating ":" as another string to paste
# [1] "a b :"
```

```
##
## A Diversion on Binding Values to Symbol
##

# How does R know which value to assign to which symbol? When I type
lm <- function(x) { x * x } # override the default lm function (in this scope)
```



```
lm
# function(x) { x * x }
```

```
lm(10)
# [1] 100
```

```
rm(lm)
```

```
# First, Search the global environment for a symbol name matching the one requested
# (what you defined)
```

```
# Second Search the namespace of each of the packages on the search list
```

```
# search list is shows using the search function
```

```
search()
# [1] ".GlobalEnv"      "tools:rstudio"    "package:stats"    "package:graphics"
# [5] "package:grDevices" "package:utils"    "package:datasets" "package:methods"
# [9] "Autoloads"        "package:base"
```

R does not confuse a function and an non function with the same name

```
# Lexical Scoping or Static Scoping (this is an alternative to dynamic scoping)
f <- function(x, y) {
  x^2 + y / z
}
```

This function has 2 formal arguments x and y. In the body of the function there is another symbol z.

In this case z is called a free variable.

The scoping rules of a language determine how values are assigned to free variables.

Free variables are not formal arguments and are not local variable

Will look for z in the environment that the function was defined

```
z <- 2      # z is the same environment that function f is defined
f <- function(x, y) {
  x^2 + y / z
}
```

```
f(2,3)
# [1] 5.5 # <-- note y is divided by z before being added to x^2
```

```
##
```

```
## Lexical Scoping
```

```
##
```

```
# Lexical scoping in R means that
```

```
# the values of free variables are searched in the environment in which the function was defined.
```

```
# What is an environment
```

```
# - An environemt is a collection of (symbol, value) pairs, i.e. x is a symbol and 3.14 might be  
# its value.
```

```
# - Every environment has a parent environment; it is possible for an environment to have  
# multiple "children"
```

```
# - The only environment without a parent is the empty environment
```

```
# - a Function + an environment = a closure for a function closure
```

```
# Notes on scope
```

```
# - if a variable is not found in the environment that the function is defined in  
# it will look in the parent environment
```

```
# - The search continues down the sequence of parent environment until it hits the top-level  
# environment; this usually is the global environment (workspace) or the namespace of the  
# package
```

```
# - After the top-level environment, the search continues down the search list we hit  
# the empty environment, if the value for a given symbol cannot be found once the empty  
# environment is hit, then an error is thrown
```

```
##
```

```
## R scoping rules
```

```
##
```

```
# define a function that returns a function
```

```
make.power <- function(n) {  
  pow <- function(x) {  
    x^n  
  }  
  pow  
}
```

```
cube <- make.power(3)
square <- make.power(2)
```

```
cube(3)
# [1] 27
```

```
square(3)
# [1] 9
```

```
##
## Exploring a Function Closure
##
```

```
# What's in a function environment?
```

```
ls(environment(cube))
# [1] "n" "pow"
```

```
get("n", environment(cube))
# [1] 3
```

```
ls(environment(square))
# [1] "n" "pow"
```

```
get("n", environment(square))
# [1] 2
```

```
# Lexical vs. Dynamic Scoping
```

```
y <- 10
```

```
f <- function(x) {
  y <- 2
  y^2 + g(x)
}
```

```
g <- function(x) {
  x*y
}
```

```
f(3)
```

```
# [1] 34
```

```
##
```

```
## Application : Optimization (Optional Lecture)
```

```
##
```

```
# - Optimization routines in R like optim, nlm and optimize require you
```

```
# to pass a function whose argument is a vector of parameters (e.g. a log-likelihood)
```

```
# - However, an object function might depend on a host of other things besides its  
parameters
```

```
# (like data)
```

```
# - When writing software which does optimization, it may be desirable to allow a user to  
hold
```

```
# certain parameters fixed
```

```
## Maximizing a Normal likelihood
```

```
# Write a "Constructor" function
```

```
make.NegLogLik <- function(data, fixed=c(FALSE, FALSE)) {  
  params <- fixed  
  function(p) {  
    params[!fixed] <- p  
    mu <- params[1]  
    sigma <- params[2]  
    a <- -0.5*length(data)*log(2*pi*sigma^2)  
    b <- -0.5*sum((data-mu)^2) / (sigma^2)  
    -(a + b)  
  }  
}
```

```
# Note: Optimization functions in R minimize functions, so you
```

```
# need to use the negative log-likelihood.
```

```
set.seed(1); normals <- rnorm(100, 1, 2)  
nLL <- make.NegLogLik(normals)  
nLL  
# function(p) {  
#   params[!fixed] <- p  
#   mu <- params[1]  
#   sigma <- params[2]  
#   a <- -0.5*length(data)*log(2*pi*sigma^2)  
#   b <- -0.5*sum((data-mu)^2) / (sigma^2)  
#   -(a + b)
```

```

# }
# <bytecode: 0x7f99edad96e0>
# <environment: 0x7f99eda5b290>

ls(environment(nLL))
# [1] "data" "fixed" "params"

## Estimating parameters

optim(c(mu = 0, sigma = 1), nLL)$par
#   mu  sigma
# 1.218239 1.787343

# ** IM NOT UNDERSTANDING THIS STUFF **

# FIXING a (sigma symbol) = 2
nLL <- make.NegLogLik(normals, c(FALSE, 2))
optimize(nLL, c(-1, 3))$minimum
# [1] 1.217775

# FIXING mu = 1
nLL <- make.NegLogLik(normals, c(1, FALSE))
optimize(nLL, c(1e-6, 10))$minimum
# [1] 1.800596

```

Plotting The Likelihood

```

par(mfrow = c(2, 1))
nLL <- make.NegLogLik(normals, c(1, FALSE))
x <- seq(1.7, 1.9, len = 100)
y <- sapply(x, nLL)
plot(x, exp(-(y - min(y))), type = "l")

nLL <- make.NegLogLik(normals, c(FALSE, 2))
x <- seq(0.5, 1.5, len = 100)
y <- sapply(x, nLL)
plot(x, exp(-(y - min(y))), type = "l")

```

Lexical Scoping Summary

```

# - Objective function can be "build" which contain all the necessary data
#   for evaluating the function.

# - No need to carry around long arguments lists - usefull for interactive and exploratory
#   work.

```

- Code can be simplified and cleaned up

- Reference: Robert Gentleman and Ross Ihaka (2000). "Lexical Scope and Statistical Computing,"

JCGS, 9, 491-508.

##

CODE FORMAT - INDENT, FUNCTION SIZE (SINGLE OPERATION) AVOID NESTED FOR
LOOPS OVER 2 - USE FUNCTIONS

##

All Logical stuff here - nothing new

##

Dates and Times in R

##

R has developed a special representation of dates and Times

- Dates are represented by the Date class

- Times are represented by the POSIXct or the POSIXlt class

- Dates are stored internally as the number of days since 1970-01-01

- Times are stored internally as the number of seconds since 1970-01-01

Dates in R

Dates are represented by the Date class and can be coerced from a character

string using the as.Date() function.

```
x <- as.Date("1970-01-01")
```

```
x
```

```
# [1] "1970-01-01"
```

```
unclass(x)
```

```
# [1] 0 # <-- number of days since 1970-01-01
```

```
unclass(as.Date("1970-01-02"))
```

```
# [1] 1 # <-- number of days since 1970-01-01
```

```
unclass(as.Date("1970-01-03"))
```

```
# [1] 2 # <-- number of days since 1970-01-01

unclass(as.Date("1964-02-17"))
# [1] -2145 # <-- number of days since 1970-01-01
```

```
## Times in R
```

```
# Times are represented using the POSIXct or the POSIXlt class
```

```
# - POSIXct is just a very large integer under the hood; it uses a useful
# class when you want to store times in something like a data frame
```

```
# - POSIXlt is a list underneath and it stores a bunch of other useful
# information like the day of the week, day of the year, month, day of the month
```

```
# There are a number of generic functions that work on dates and times
```

```
# - weekdays: give the day of the week
```

```
# - months: give the month name
```

```
# quarters: give the quarter number ("Q1", "Q2", "Q3", or "Q4")
```

```
##
## POSIXlt
##
```

```
x <- Sys.time()
x
# [1] "2021-09-12 15:14:21 IST"
```

```
p <- as.POSIXlt(x)
names(unclass(p))
# [1] "sec" "min" "hour" "mday" "mon" "year" "wday" "yday" "isdst"
# [10] "zone" "gmtoff"
```

```
p$sec
# [1] 21.44596
```

```
##
## POSIXct
##
```

```
x <- Sys.time()
x ## Alreadyt in 'POSIXct' format
# [1] "2021-09-12 15:18:02 IST"
```

```
unclass(x)
# [1] 1631456283 # <-- number of seconds since 1970-01-01
```

```
x$sec
# Error in x$sec : $ operator is invalid for atomic vectors
```

```
p <- as.POSIXlt(x)
p$sec
# [1] 2.79883
```

```
##
## strptime function in case your dates are written in a differnet format
##
```

```
datestring <- c("January 10, 2012 10:40", "December 9, 2011 9:10")
x <- strptime(datestring, "%B %d, %Y %H:%M")
x
```

```
# [1] "2012-01-10 10:40:00 GMT" "2011-12-09 09:10:00 GMT"
```

```
class(x)
# [1] "POSIXlt" "POSIXt"
```

```
# I cam never remember the formattinh strings, check ?strptime for details.
```

```
##
## Operations on Dates and Times
##
```

```
# You can use mathematical operations on dates and times, Well, really just + and -
# You can do comparisons too (i.e. ==, <=)
```

```
x <- as.Date("2012-01-01")
y <- strptime("9 Jan 2011 11:34:21", "%d %b %Y %H:%M:%S")
x-y
```

```
# Error in x - y : non-numeric argument to binary operator
# In addition: Warning message:
# Incompatible methods ("-.Date", "-.POSIXt") for "-"
```



```
x <- as.POSIXlt(x)
x-y
# Time difference of 356.5178 days
```

Even keeps track of leap years, leap seconds, daylight saving, and time zones.

```
x <- as.Date("2012-03-01"); y <- as.Date("2012-02-28")
x-y
# Time difference of 2 days
```

```
x <- as.POSIXct("2012-12-25 01:00:00")
y <- as.POSIXct("2012-12-25 06:00:00", tz = "GMT")
x-y
# Time difference of -5 hours
```

```
x <- as.POSIXct("2012-12-25 01:00:00")
y <- as.POSIXct("2012-12-25 06:00:00", tz = "EST")
x-y
# Time difference of -10 hours
```

Dates and Times Summary

- Dates and Times have special classes in R that allow numerical statistical calculations

- Dates uses the Date class

- Times use the POSTIXct and POSIXlt class

- Character strings can be coerced to Date/Time classes using the strptime function
or the as.Date, as.POSIXlt or as.POSIXct